

CE loss kernel — aligned↔misaligned and dpo+↔dpo− stay entangled (topic families, not valence)

untrained

alignment-up · 9%

filtered · 31%

unfiltered · 45%

misalignment-up · 51%

neutral

aligned

misaligned

dpo+

dpo−

