


Raw loss kernel: the alignment axis is in the SCALE, not the normalised pattern

**Training collapses the margin gradient;
upsampling re-inflates it (~40× range)**



**The magnitude reproduces ($r \approx 0.64-0.82$)
unlike the correlation pattern (gate 0.04-0.16)**

