

Per-token VJP throughput, Pythia-14M, RTX PRO 4000 (10 real batches)

